

Predicting Solar Flares with Space Weather Data

1. Introduction:

Purpose of the Project:

The purpose of this project is to build a predictive model for time series forecasting using solar wind and sunspot data. The model will utilize the **LSTM** (Long Short-Term Memory) neural network, which is well-suited for handling sequential data and capturing long-term dependencies.

Project Overview:

The project involves several steps, including data preprocessing, exploratory data analysis, building the LSTM model, training the model on the data, evaluating its performance, and finally saving the trained model for future use.

Required Libraries:

The following Python libraries and modules are utilized throughout the project:

- NumPy - For numerical operations and array manipulation.
- Pandas - For data manipulation and analysis.
- Matplotlib - For data visualization and plotting.
- Seaborn - For enhanced data visualization.
- SciPy - For statistical and scientific computations.
- Scikit-learn - For machine learning tasks and preprocessing.
- TensorFlow - For building and training the LSTM model.
- Keras - As a high-level API for building neural networks.

2. Data Description:

Solar Wind Data:

- Description: The solar wind data contains measurements of various solar wind parameters, including magnetic field components (bx_gse, by_gse, bz_gse), temperature, speed, and density.
- Format: The data is structured as a CSV file with rows representing different time periods and columns for each parameter.
- Time Periods: The data is collected at regular intervals over different periods.

Sunspot Data:

- Description: The sunspot data provides information about the number of sunspots observed at different times.
- Format: The data is stored in a CSV file with rows representing different time periods and a column for the sunspot count.

DST Labels Data:

- Description: The DST labels data contains the disturbance storm time index (DST) values corresponding to different time periods.
- Format: The data is available in a CSV file with rows representing different time periods and a column for the DST values.

3. Data Preprocessing:

Data Loading:

- The script loads the solar wind, sunspot, and DST labels data from their respective CSV files into Pandas DataFrames.
- Time-related columns are converted to the proper datetime format using Pandas' "to_timedelta" function.
- The "period" and "timedelta" columns are used to set a multi-level index for each DataFrame.

Data Cleaning:

- The script checks for missing values in each DataFrame.
- Any missing values in the sunspot data are filled using forward filling (method="ffill").
- For the solar wind data, missing values are interpolated using the DataFrame's "interpolate" function.

Data Joining:

- The solar wind, sunspot, and DST labels data are joined based on their index (period and timedelta) using Pandas' "join" function.
- The join operation ensures that the data from different sources is aligned for each time period.

Feature Selection:

A subset of solar wind features is selected for modeling. The chosen features are "bt," "temperature," "bx_gse," "by_gse," "bz_gse," "speed," and "density."

Feature Scaling:

- The selected solar wind features are scaled using standard scaling to bring them to a similar range.
- StandardScaler from scikit-learn is used for this purpose.

Missing Value Imputation:

- Any remaining missing values in the joined DataFrame are imputed using forward filling for sunspot data and interpolation for solar wind data.

4. Exploratory Data Analysis (EDA):

Visualizing Raw Solar Wind Data:

- The script defines a function "show_raw_visualization" to visualize the raw solar wind data for the selected features.
- The first 1000 rows of the selected columns are plotted using Matplotlib to provide an initial view of the data's characteristics.

Correlation Analysis:

- Correlation matrices are generated for the features to understand the relationships between them.
- The correlations are visualized using Seaborn's "clustermap" function to identify any strong associations between variables.

5. Model Building and Training:

LSTM Neural Network Architecture:

- The LSTM model is chosen for time series forecasting due to its ability to capture long-term dependencies in sequential data.
- The model is built using the Keras Sequential API.

Data Preparation for LSTM:

- The script defines a function "timeseries_dataset_from_df" to prepare the data in time series format for LSTM.
- The input sequences and corresponding target values are created from the DataFrame.

Model Compilation:

- The LSTM model is compiled with the mean squared error loss function ("mean_squared_error") and the Adam optimizer.

Model Training:

- The LSTM model is trained on the training data for a specified number of epochs using the "fit" function.
- The training process includes backpropagation to update the model's weights and biases.

Model Evaluation:

- The trained LSTM model is evaluated on the validation dataset to assess its performance.
- The model's loss and validation metrics are recorded and can be visualized to monitor training progress.

6. Model Testing:

Test Dataset Preparation:

- The test dataset is prepared to evaluate the model's performance on unseen data.
- Similar to the validation data preparation, the test data is formatted into input sequences and target values.

Model Evaluation on Test Dataset:

- The trained LSTM model is evaluated on the test dataset to calculate the root mean squared error (RMSE) as a performance metric.
- RMSE is a measure of the difference between the predicted and actual values.

Performance Metrics:

- The model's performance metrics, including the test RMSE value, are printed for evaluation.

7. Model Saving and Configuration:

Saving Trained Model:

- The trained LSTM model is saved using Keras' "model.save" function.
- The saved model can be loaded and used for predictions in future applications.

Saving Scaler Object:

- The fitted StandardScaler object, used for scaling the features during training, is saved using Python's "pickle" module.
- The saved scaler can be reused to transform new data consistently.

Saving Configuration as JSON:

- The project configuration data, including data parameters and feature selections, is saved in a JSON file for reference.
- The configuration information can be used for reproducing the same data preprocessing and model settings in the future.

8. Conclusion:

Summary of the Project:

- The project successfully implements time series forecasting using LSTM neural networks with solar wind and sunspot data.
- It demonstrates the data preprocessing, modeling, and evaluation steps involved in building a predictive model for time series data.

Key Findings:

The LSTM model exhibits promising performance in forecasting DST labels, as indicated by the evaluation metrics.

Future Enhancements:

The documentation suggests potential improvements or future enhancements that can be explored to enhance the model's accuracy and robustness.

This detailed documentation provides a step-by-step explanation of the entire project, covering data loading, preprocessing, exploratory data analysis, LSTM model building, training, evaluation, and model saving. It serves as a comprehensive reference guide for understanding and replicating the time series forecasting project.

Certainly! Let's explain the approach and code step by step:

1. Approach:

The code's main objective is to build a predictive model for time series forecasting using solar wind and sunspot data. It follows these key steps:

- **Data Loading:** The script reads three CSV files containing solar wind data, sunspot data, and labels (dst - disturbance storm time index) data.
- **Data Preprocessing:** The data is preprocessed, which includes converting time-related columns to proper datetime format, setting multi-level indices, handling missing values, and joining the dataframes.
- **Feature Selection and Scaling:** A subset of solar wind features is selected for modeling. The features are aggregated to hourly intervals, and standard scaling is applied to the data.
- **Train-Test-Validation Split:** The data is split into training, testing, and validation sets. This ensures the model is tested on unseen data to assess its generalization performance.

- **Model Building:** An LSTM neural network is chosen for time series forecasting due to its ability to capture long-term dependencies in sequential data.
- **Model Training and Evaluation:** The LSTM model is trained on the training data and evaluated on the validation and test datasets.
- **Model Saving:** The trained LSTM model is saved for future use along with the configuration and scaling information.

2. Code Explanation:

The code starts by importing necessary libraries and modules for data manipulation, visualization, machine learning, and deep learning.

- Next, it loads the three CSV files ("solar_wind.csv," "labels.csv," and "sunspots.csv") into Pandas DataFrames ("solar_wind," "dst," and "sunspots").
- The time-related columns are converted to Pandas Timedelta format, and the multi-level index is set using the columns "period" and "timedelta" for each DataFrame.
- The script defines a function "show_raw_visualization" to visualize the first 1000 rows of selected columns from the "solar_wind" DataFrame using Matplotlib.
- Data imputation is performed for sunspot data using forward filling and for solar wind data using interpolation.
- The script aggregates solar wind features to the floor of each hour using mean and standard deviation.
- Feature preprocessing involves selecting the desired solar wind features, aggregating, joining with sunspots, scaling using standard scaling, and imputing missing values.
- A function "get_train_test_val" splits the data into train, test, and validation sets based on the number of rows to be assigned for each period.
- The LSTM model is defined using the Keras Sequential API. It consists of an LSTM layer with a specified number of neurons and a dense output layer with the number of output features.
- The model is compiled with the mean squared error loss function and the Adam optimizer.
- The training data is prepared as time series datasets using the "timeseries_dataset_from_df" function, and the model is trained for a specified number of epochs.

- The model's training history is stored to visualize the loss and validation metrics over epochs.
- The model's performance is evaluated on the test dataset using the mean squared error metric.
- Finally, the trained LSTM model is saved using Keras' "model.save" function, and the configuration data is saved as JSON, along with the fitted scaler for future use.

In summary, the code implements a comprehensive workflow for time series forecasting using LSTM neural networks, including data loading, preprocessing, model building, training, evaluation, and model saving.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from scipy import stats
from scipy.stats import norm, skew

from sklearn.model_selection import train_test_split, KFold, GroupKFold, GridSearchCV,
StratifiedKFold
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler,
MinMaxScaler, RobustScaler

from sklearn.metrics import *

import sys, os
import random

if not sys.warnoptions:
    import warnings
    warnings.simplefilter("ignore")

# Read the data files: solar wind, DST labels, and sunspots
solar_wind = pd.read_csv("Downloads/solar_wind.csv")
dst = pd.read_csv("Downloads/labels.csv")
sunspots = pd.read_csv("Downloads/sunspots.csv")

# Convert timedelta column to proper datetime format
solar_wind.timedelta = pd.to_timedelta(solar_wind.timedelta)
dst.timedelta = pd.to_timedelta(dst.timedelta)
sunspots.timedelta = pd.to_timedelta(sunspots.timedelta)

# Set the index for each DataFrame using "period" and "timedelta"
solar_wind.set_index(["period", "timedelta"], inplace=True)
```

```

dst.set_index(["period", "timedelta"], inplace=True)
sunspots.set_index(["period", "timedelta"], inplace=True)

# Print basic information about the data and show the first few rows
print("Solar wind shape: ", solar_wind.shape)
print("Sunspot shape: ", sunspots.shape)
print("Solar wind data:")
print(solar_wind.head())
print("Sunspot data:")
print(sunspots.head())

# Visualize the raw solar wind data for selected features
plt.style.use('fivethirtyeight')

def show_raw_visualization(data):
    fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(15, 15), dpi=80)
    for i, key in enumerate(data.columns):
        t_data = data[key]
        ax = t_data.plot(
            ax=axes[i // 2, i % 2],
            title=f"{key.capitalize()}",
            rot=25, color='teal', lw=1.2
        )

    fig.subplots_adjust(hspace=0.8)
    plt.tight_layout()

cols_to_plot = ["bx_gse", "bx_gsm", "bt", "density", "speed", "temperature"]
show_raw_visualization(solar_wind[cols_to_plot].iloc[:1000])

# Check for missing values in solar wind data
print(solar_wind.isnull().sum())

# Join the solar wind, sunspot, and DST labels data and fill missing values
joined = solar_wind.join(sunspots).join(dst).fillna(method="ffill")

# Visualize the correlation between features using a clustermap
plt.figure(figsize=(20, 15))
sns.clustermap(joined.corr(), annot=True)

# Set random seeds for reproducibility
from numpy.random import seed
from tensorflow.random import set_seed
seed(2020)
set_seed(2021)

# Import StandardScaler for feature scaling
from sklearn.preprocessing import StandardScaler

```



```

# Define a subset of solar wind features to use for modeling
SOLAR_WIND_FEATURES = [
    "bt",
    "temperature",
    "bx_gse",
    "by_gse",
    "bz_gse",
    "speed",
    "density",
]

# Define all features to use, including sunspot numbers
XCOLS = (
    [col + "_mean" for col in SOLAR_WIND_FEATURES]
    + [col + "_std" for col in SOLAR_WIND_FEATURES]
    + ["smoothed_ssn"]
)

# Function to impute missing values using forward fill for sunspot data and interpolation for
solar wind data
def impute_features(feature_df):
    feature_df.smoothed_ssn = feature_df.smoothed_ssn.fillna(method="ffill")
    feature_df = feature_df.interpolate()
    return feature_df

# Function to aggregate features to the floor of each hour using mean and standard
deviation
def aggregate_hourly(feature_df, aggs=["mean", "std"]):
    agged = feature_df.groupby(
        ["period", feature_df.index.get_level_values(1).floor("H")]
    ).agg(aggs)
    agged.columns = ["_".join(x) for x in agged.columns]
    return agged

# Function to preprocess features
def preprocess_features(solar_wind, sunspots, scaler=None, subset=None):
    if subset:
        solar_wind = solar_wind[subset]

    hourly_features = aggregate_hourly(solar_wind).join(sunspots)

    if scaler is None:
        scaler = StandardScaler()
        scaler.fit(hourly_features)

    normalized = pd.DataFrame(
        scaler.transform(hourly_features),

```

```

        index=hourly_features.index,
        columns=hourly_features.columns,
    )

    imputed = impute_features(normalized)

    return imputed, scaler

# Preprocess features using the defined functions
features, scaler = preprocess_features(solar_wind, sunspots,
subset=SOLAR_WIND_FEATURES)

# Check the shape and ensure there are no missing values in the features DataFrame
print(features.shape)
assert (features.isna().sum() == 0).all()

# Define target columns for labels
YCOLS = ["t0", "t1"]

# Function to process labels for forecasting
def process_labels(dst):
    y = dst.copy()
    y["t1"] = y.groupby("period").dst.shift(-1)
    y.columns = YCOLS
    return y

# Process labels using the defined function
labels = process_labels(dst)

# Combine features and labels into a single DataFrame
data = labels.join(features)

# Function to split data across periods into train, test, and validation sets
def get_train_test_val(data, test_per_period, val_per_period):
    test = data.groupby("period").tail(test_per_period)
    interim = data[~data.index.isin(test.index)]
    val = data.groupby("period").tail(val_per_period)
    train = interim[~interim.index.isin(val.index)]
    return train, test, val

# Split data into train, test, and validation sets using the defined function
train, test, val = get_train_test_val(data, test_per_period=6_000, val_per_period=3_000)

import tensorflow as tf
from keras import preprocessing

# Configuration for time series data
data_config = {

```

```

    "timesteps": 32,
    "batch_size": 32,
}

```

```

# Function to create time series dataset from DataFrame

```

```

def timeseries_dataset_from_df(df, batch_size):
    dataset = None
    timesteps = data_config["timesteps"]

    for _, period_df in df.groupby("period"):
        inputs = period_df[XCOLS][-timesteps:]
        outputs = period_df[YCOLS][timesteps:]
        period_ds = tf.keras.preprocessing.timeseries_dataset_from_array(
            inputs,
            outputs,
            timesteps,
            batch_size=batch_size,
        )

        if dataset is None:
            dataset = period_ds
        else:
            dataset = dataset.concatenate(period_ds)

    return dataset

```

```

# Create time series datasets for training and validation

```

```

train_ds = timeseries_dataset_from_df(train, data_config["batch_size"])
val_ds = timeseries_dataset_from_df(val, data_config["batch_size"])

```

```

print(f"Number of train batches: {len(train_ds)}")
print(f"Number of val batches: {len(val_ds)}")

```

```

from tensorflow.keras import Sequential
from tensorflow.keras.layers import Conv2D, Flatten, Dense, LSTM

```

```

# Define the LSTM model

```

```

model_config = {"n_epochs": 20, "n_neurons": 512, "dropout": 0.4, "stateful": False}

```

```

model = Sequential()
model.add(
    LSTM(
        model_config["n_neurons"],
        batch_input_shape=(None, data_config["timesteps"], len(XCOLS)),
        stateful=model_config["stateful"],
        dropout=model_config["dropout"],
    )
)

```

```

model.add(Dense(len(YCOLS)))
model.compile(
    loss="mean_squared_error",
    optimizer="adam",
)

model.summary()

# Train the model and save training history
history = model.fit(
    train_ds,
    batch_size=data_config["batch_size"],
    epochs=model_config["n_epochs"],
    verbose=1,
    shuffle=False,
    validation_data=val_ds,
)

# Plot the training and validation loss
for name, values in history.history.items():
    plt.plot(values)

# Create the time series dataset for testing
test_ds = timeseries_dataset_from_df(test, data_config["batch_size"])

# Evaluate the model on the test dataset and calculate the RMSE
mse = model.evaluate(test_ds)
print(f"Test RMSE: {mse**.5:.2f}")

# Save the trained model, scaler, and data configuration for future use
model.save("model")

with open("scaler.pck", "wb") as f:
    pickle.dump(scaler, f)

data_config["solar_wind_subset"] = SOLAR_WIND_FEATURES
print(data_config)

with open("config.json", "w") as f:
    json.dump(data_config, f)

```

The code implements the entire process of time series forecasting using LSTM. It loads solar wind, sunspot, and DST labels data, performs data preprocessing, creates time series datasets, builds an LSTM model, trains the model, evaluates its performance, and finally saves the model and other necessary configurations for future use. The code is thoroughly commented to explain each step and ensure ease of understanding.